

Support Vector Machines: An Introduction

Seminar Data Mining

Chengjie "Jay" Zhou

Department of Informatics

Technische Universität München

Email: chengjie.zhou@mytum.de

Abstract—Text of the abstract

Keywords—Data Mining, Support Vector Machines

I. INTRODUCTION

One of the most important aspects of machine learning is classification. Typical algorithms and models that are used for classification include logistic regression, naive bayes, decision trees and so on. In the early 90s, Vladimir Vapnik and his colleagues developed a new algorithm called *Support Vector Machine (SVM)*, which is optimized for classification and regression analysis. In this paper, we will focus on the main usages of SVM, including the generalization of linear decision boundaries for classification. We will also discuss the role of kernel in SVM, as well as the implementation, evaluation methods, application and many different aspects of SVM.

In order to thoroughly understand the classification problem, we first need to look at a simple example in one dimension. Suppose there are two groups of data points that are separately distributed on a one-dimension number line. The classification then becomes obvious: The threshold that separates both groups simply lies in the middle of the the most outward data points from both groups. This classification method is called the *maximum margin classifier*, since the margin that separates both groups is maximized.




Fig. 1. Maximum Margin Classifier

Nevertheless, the maximum margin classifier is not always applicable, since the data points are not always ideally distributed in a separated way. If an outlier happens to appear in the dataset, it would push the threshold to one side, since the data point is much closer to the other group. This would result in a severe misclassification, since the data that are close to one group now belongs to the other group because of the shift of the threshold. A solution to this problem is to allow

some misclassification, so that the threshold has higher bias and is less sensitive to outliers and the classifier performs better when there is new data. This margin that allows some misclassification is called the *soft margin*. The determination of soft margin could be tricky, since there are limitless points of thresholds to choose from. One way to find the optimal threshold is to use Cross Validation. Cross Validation is a method that splits the dataset into several parts. For each repetition, one part of the dataset is used for testing and the rest is used for training. After training through all the repetition, the average position of the threshold represents the most ideal position of the soft margin. This classifier is called soft margin classifier, also known as support vector classifier. The name "Support Vector" derives from the fact that the data points that are closest to the threshold are called support vectors.




Fig. 2. Support Vector Classifier

In reality, chances are that the data is not only one-dimensional. In fact, almost all of the data that we work with is multi-dimensional. In order to represent the data in a multi-dimensional space, we need to use vectors. In this case, the idea of hyperplane is introduced. A *hyperplane* is a threshold that separates the data in a multi-dimensional space, which is formally defined as a flat affine subspace. [1] The support vector classifier is able to deal with this case as well. It finds the hyperplane that separates the data in a way that the margin is maximized. The exact algorithms and the mathematical derivation will be discussed in the next section.

However, the support vector classifier is still not suitable for data that is not linearly separable, even though the support vector classifiers allows misclassifications and is less sensitive to outliers.

Usually, we use a *kernel* to deal with this case, which is a function that maps the data into a higher dimensional




Fig. 3. Non Linearly Separable Data

space, so that the data becomes linearly separable. The SVM implements the idea of kernel without transforming the data into a higher dimensional space. This concept is called *the kernel trick* and is a crucial part of SVM. This trick allows SVM to cast nonlinear variants to dot products, enabling easier computation and better performance. [2]




Fig. 4. Kernel Transformation

II. IMPLEMENTATION

The problem of support vector machine is an optimization problem. The goal is, as mentioned, to find the optimal hyperplane that separates the data in a way that the margin is maximized. Therefore, it is important to first define the hyperplane.

Since the hyperplane is a flat affine subspace of dimension p , we can easily define a hyperplane as [1]:

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0 \quad (1)$$

We can then denote X and β as vectors [3]:

$$\{x \in \mathbb{R}^p : X^T \beta + \beta_0 = 0\} \quad (2)$$

With hyperplane being defined, we are now able to observe the problem of separating the hyperplane. In other words, we need to construct linear decision boundaries that attempt to separate the data into two or more classes as precisely as possible. The main two mechanics to this problem are Rosenblatt and Optimal.

The *Rosenblatt's perceptron learning algorithm* aims to find the optimal hyperplane in an iterative way. The algorithm minimizes the distance of misclassified points to the decision

boundary. The goal of this algorithm is to minimize the following equation

$$D(\beta, \beta_0) = \sum_{i \in M} -y_i(x_i^T \beta + \beta_0) \quad (3)$$

where M is the set of misclassified points and y_i is the class label of x_i , being either 1 or -1. Using stochastic gradient descent, the algorithm updates the coefficients and intercept.

$$\frac{\partial D(\beta, \beta_0)}{\partial \beta} = - \sum_{i \in M} y_i x_i \quad (4)$$

$$\frac{\partial D(\beta, \beta_0)}{\partial \beta_0} = - \sum_{i \in M} y_i \quad (5)$$

However, the perceptron learning algorithm is not deterministic, since the result depends on the starting values. Furthermore, the algorithm does not converge if the data is not linearly separable, resulting in an infinite loop if the case was not detected beforehand. [4]

In 1996, Vapnik proposed the *optimal separating hyperplane* algorithm in order to cope with the problems of the perceptron learning algorithm. This algorithm is widely implemented as maximum margin classifier. The goal of the algorithm is to find a hyperplane that maximize the distance to the closet point from either class. [Vapnik, 1996] The observation of both classes is denoted as $y_i = 1$ and $y_i = -1$. The separating hyperplane then has the following property:

$$X^T \beta + \beta_0 \geq 1, \text{ if } y_i = 1 \quad (6)$$

$$X^T \beta + \beta_0 \leq -1, \text{ if } y_i = -1 \quad (7)$$

To summarize:

$$y_i(X^T \beta + \beta_0) \geq 1, \quad i = 1, \dots, n \quad (8)$$

To solve this problem, we construct the following optimization problem [1]:

$$\begin{aligned} & \text{maximize} \quad M \\ & \beta, \beta_0 \\ & \text{subject to} \\ & \|\beta\| = 1, \\ & y_i(x_i^T \beta + \beta_0) \geq M, \quad i = 1, \dots, n \end{aligned} \quad (9)$$

where M is the margin. Since the margin is M units away from the hyperplane on either side, the margin is then $2M$ units wide. The constraint $\|\beta\| = 1$ can be also left out by replacing the condition by

$$\frac{1}{\|\beta\|} y_i(x_i^T \beta + \beta_0) \geq M \quad (10)$$

We arbitrarily set $M = \frac{1}{\|\beta\|}$. The problem can be then reconstructed as [4]:

$$\begin{aligned} & \text{maximize} \quad \frac{1}{2} \|\beta\|^2 \\ & \beta, \beta_0 \\ & \text{subject to} \\ & y_i(x_i^T \beta + \beta_0) \geq 1, \quad i = 1, \dots, n \end{aligned} \quad (11)$$

The two constraints of this optimization problem ensures that each observation is on the right side and has at least a distance of M from the hyperplane. It can be then solved in an efficient way using Lagrange functions. The mathematical deduction is beyond the scope of this paper, but the result is as follows:

$$\beta = \sum_{i=1}^n \alpha_i y_i x_i \quad (12)$$

$$0 = \sum_{i=1}^n \alpha_i y_i \quad (13)$$

$$\alpha_i [y_i (x_i^T \beta + \beta_0) - 1] = 0 \quad \forall i \quad (14)$$

where α is a vector of the weights of all the training points as support vectors under the condition of $\alpha_i \geq 0$, $i = 1, \dots, n$. From these three equations, we can derive the following two properties [4]:

- if $\alpha_i > 0$, then $y_i (x_i^T \beta + \beta_0) - 1 = 0$, which is equivalent to $y_i (x_i^T \beta + \beta_0) = 1$. This means that the observation is on the margin.
- if $y_i (x_i^T \beta + \beta_0) - 1 > 0$, or $y_i (x_i^T \beta + \beta_0) > 1$, then $\alpha_i = 0$. This means that the observation is not on the margin.

However, the maximum margin classifier has zero tolerance for misclassification. As mentioned before, the support vector is designed to be more susceptible to misclassification. With the mathematical knowledge given above, we can also define the support vector classifier. The hyperplane is chosen to correctly separate most of the dataset into two classes with the tolerance of a few errors. Thus we introduce a slack variable ϵ_i , which allows individual observations to be on the wrong side of the margin. The optimization problem is then defined as [1]:

$$\begin{aligned} & \text{maximize} \quad M \\ & \beta, \beta_0, \epsilon, M \\ & \text{subject to} \\ & ||\beta|| = 1, \\ & y_i (x_i^T \beta + \beta_0) \geq M(1 - \epsilon_i), \quad i = 1, \dots, n, \\ & \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C \end{aligned} \quad (15)$$

The interpretation of the slack variable helps us to determine whether an observation is on the correct side of the margin. The i th observation is on the correct side of the margin if $\epsilon_i = 0$. Otherwise, the i th observation is on the wrong side of the margin if $\epsilon_i > 0$ and on the wrong side of the hyperplane if $\epsilon_i > 1$. The parameter C , in this case, is the tuning paramter. It determines the budget of violations that could happen in this classification problem. The choice of C plays a decisive role in the learning process. If C is large, then many observations violate the margin, which causes the involvement of many observations in determining the hyperplane. On the other hand, smaller C s result in a classifier with lower bias but higher variance. In practice, C is chosen using cross-validation.

Based on the optimization problem given above, we can derive the formal definition of SVM. In comparison to support vector classifier, the SVM transform the data into a higher dimension space p' , so that the data could be linearly separable. For instance, we transform the features X_1, X_2, \dots, X_p into $X_1, X_1^{p'}, X_2^{p'}, \dots, X_p^{p'}$. The optimization problem with the transformed data can then be altered to [1]:

$$\begin{aligned} & \text{maximize} \quad M \\ & \beta_0, \beta_{11}, \beta_{12}, \dots, \beta_{p(p-1)'}, \beta_{pp'}, \epsilon_1, \dots, \epsilon_n, M \\ & \text{subject to} \\ & y_1 (\beta_0 + \sum_{j=1}^p \sum_{k=1}^{p'} \beta_{jk} x_{ij}^k) \geq M(1 - \epsilon_i) \\ & \sum_{i=1}^n \epsilon_i \leq C, \quad \epsilon_i \geq 0, \quad \sum_{j=1}^p \sum_{k=1}^{p'} \beta_{jk}^2 = 1 \end{aligned} \quad (16)$$

In order to perform the transformation and simplify this problem, we use a kernel function $K(x, x')$. It is defined as [3]:

$$K(x, x') = \langle h(x), h(x') \rangle \quad (17)$$

where $h(x)$ is a transformation function. The transformation process leaves us with a lot of freedom, and we are able to freely decide how we should transform the data. The two popular choices of kernel functions in practical applications are:

- dth-Degree Polynomial kernel: $K(x, x') = (1 + \langle x, x' \rangle)^d$
- Radial kernel: $K(x, x') = \exp(-\gamma ||x - x'||^2)$

Let us first take a look at the polynomial kernel. The polynomial kernel transforms each data point into a polynomial of degree d and fits a support vector classifier into this higher-dimensional space. This transformation leads to a flexible decision boundary.

$$K(x, x') = (1 + \langle x, x' \rangle)^d = (1 + \sum_{j=1}^p x_{ij} x_{i'j})^d \quad (18)$$

Another popular choice is the radial kernel. In order to understand the radial kernel, we first need to understand the following scenario. If a given test observation x^* is far away from the observation x_i , then the Euclidean distance between the two points will be large. As a result, the exponential term of the negative value of the Euclidean distance will be close to zero. In this case, the presence of x^* will not have much effect on the transformation. On the other hand, if x^* is close to x_i , then the exponential term will be close to one, which will exert a strong influence on the transformation. The radial kernel utilizes this property of the exponential function to emphasize the points that are close to the test observation x^* [1].

$$K(x, x') = \exp(-\gamma ||x - x'||^2) = \exp(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2) \quad (19)$$

The implementation of kernels in SVM saves computational power, since for each pair of data points, we only need to calculate the transformed value once. Furthermore, the transformative property of the kernel leads to a flexible and precise boundary. [1]

III. CHAPTER-1

A. Subchapter

blabla with three references [?], [?], [?]




Fig. 5. Tree

TABLE I
BEISPIELTABELLE

Spalte1	Spalte2
0	1

IV. SUMMARY AND OUTLOOK

blabla

REFERENCES

[1] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning with Applications in R*, second edition ed. New York: Springer, 2021, ch. 9, Support Vector Machines.

[2] T. Hofmann, B. Schölkopf, and A. J. Smola, *Learning with Kernels: Support Vector Machines, Regularization, Optimization, and Beyond*. Cambridge, MA: The MIT Press, 2002, ch. 2, Kernels.

[3] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, second edition ed. New York: Springer, 2009, ch. 12, Support Vector Machines and Flexible Discriminants.

[4] —, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, second edition ed. New York: Springer, 2009, ch. 4, Linear Methods for Classification.